



# Big data platform (Spark) performance acceleration

*Mentors: Tony Tan, Ning Wu, Yong Wang and Theo Gkountouvas*

By:

Grishma Atul Thakkar

Virat Goradia

Nipun Midha

Baoshu Brady Qi



# Sprint Goals


- Further understand existing spark internal code
- Design strategies to implement N-Way merge algorithm
- Start implementing the N-Way merge algorithm.

# Burndown Chart

DEMO 3 NEU6620-BIG DATA PLATFORM (SP... 18 OCT 2019-25 OCT 2019




 70%  $\vee$  80 total points 56 completed points | 0 open tasks 0 closed tasks  $\rightleftharpoons$  |  0 iocaine doses




# Spark Codebase

839,842 lines of code



## STEP BY STEP DEBUGGING

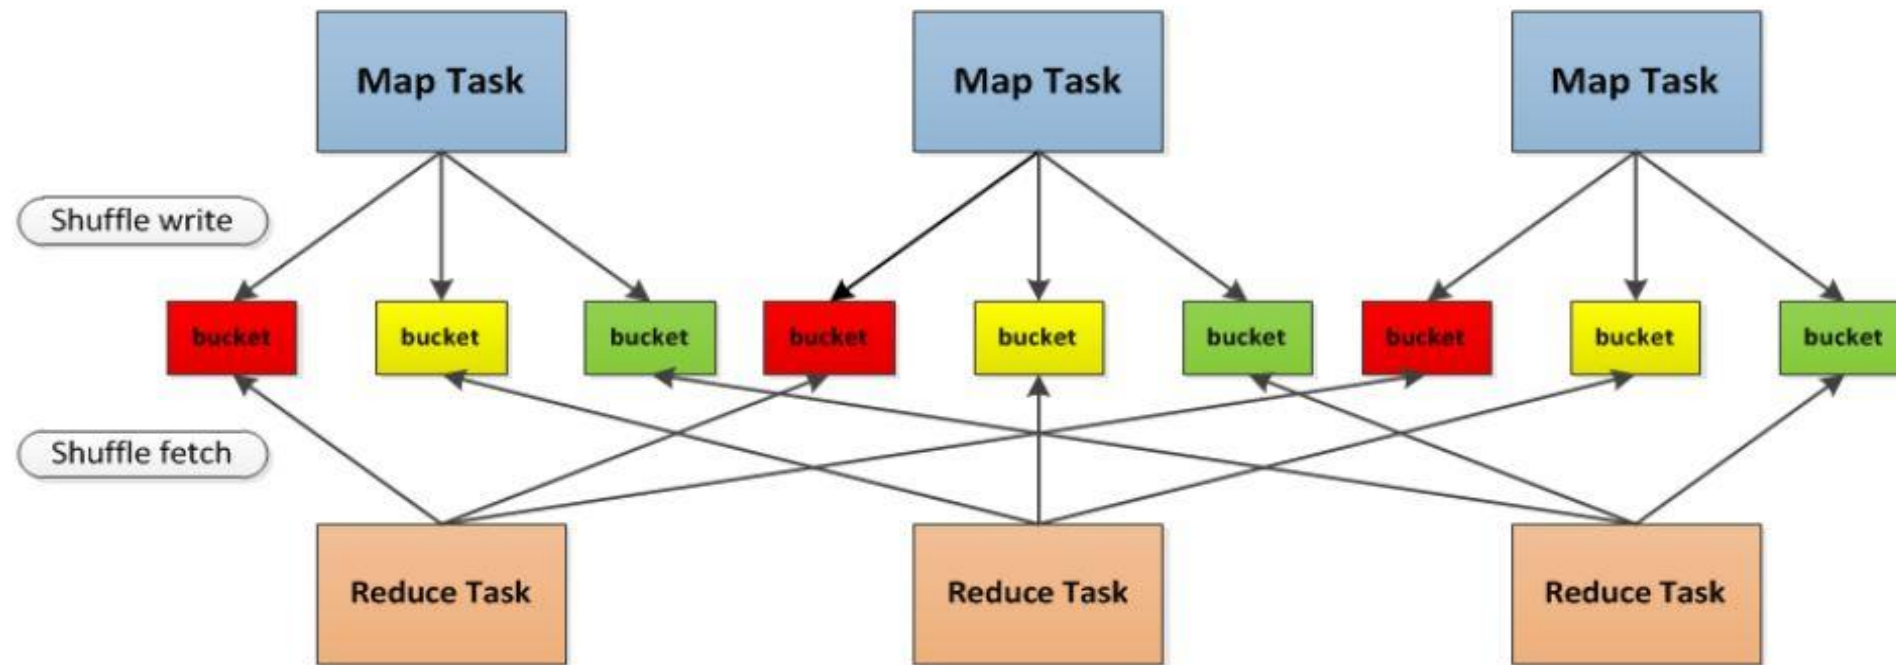




# Challenges

---

# Spark's Shuffle Phase



# ShuffleWrite - Older Version

---

```
override def run(attemptId: Long): MapStatus = {
  val numOutputSplits = dep.partitioner.numPartitions

  ...
  // Partition the map output.
  val buckets = Array.fill(numOutputSplits)(new ArrayBuffer[(Any, Any)])
  for (elem <- rdd.iterator(split, taskContext)) {
    val pair = elem.asInstanceOf[(Any, Any)]
    val bucketId = dep.partitioner.getPartition(pair._1)
    buckets(bucketId) += pair
  }

  ...

  val blockManager = SparkEnv.get.blockManager
  for (i <- 0 until numOutputSplits) {
    val blockId = "shuffle_" + dep.shuffleId + "_" + partition + "_" + i
    // Get a Scala iterator from Java map
    val iter: Iterator[(Any, Any)] = buckets(i).iterator
    val size = blockManager.put(blockId, iter, StorageLevel.DISK_ONLY, false)
    totalBytes += size
  }
  ...
}
```

# Problems with it

---

1. Spark will first save all the data in memory and then dump it into disk

2. Each mapper will generate a small file for each reducer. I/O will be significantly slowed when transfer large amount of small pieces of files.



# ShuffleBlockManager

---

Spark then introduces ShuffleBlockManager, which assigns a DiskObjectWriter to each bucket, to manage the shuffling files .

```
// Write the map output to its associated buckets.
for (elem <- rdd.iterator(split, taskContext)) {
    val pair = elem.asInstanceOf[Product2[Any, Any]]
    val bucketId = dep.partitioner.getPartition(pair._1)
    buckets.writers(bucketId).write(pair)
}
```

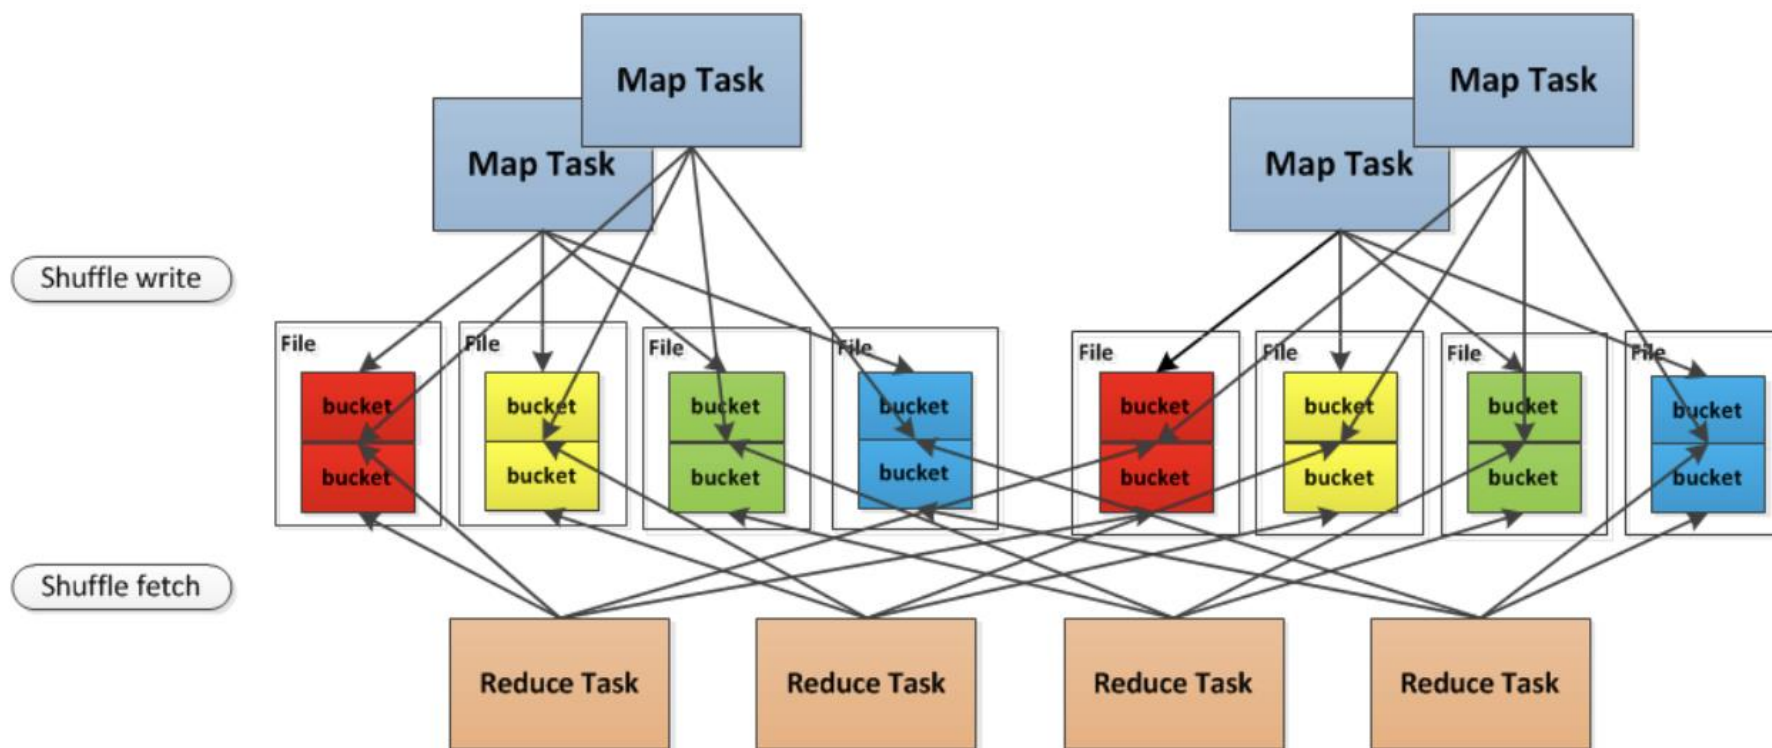
# Problems still exist

---

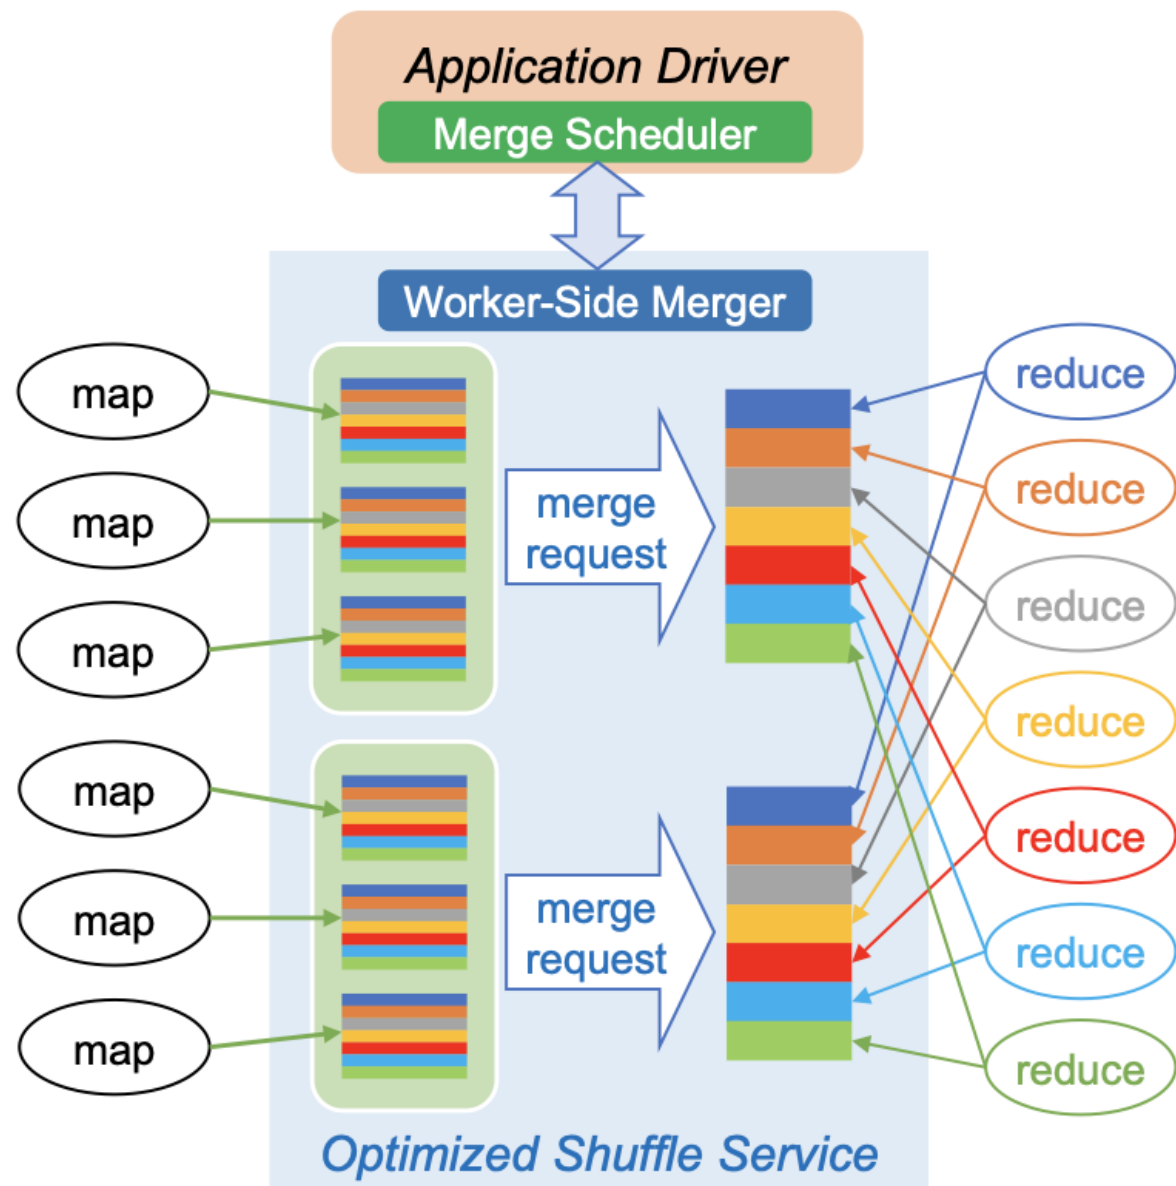
Problem with large amount of small files still exists.

Number of fix-sized writer buffer depends on the number of reducers. For a 8 core machine, 1k reducer requires 800 MB memory.


# Shuffle consolidation



Go back to  
Riffle paper



## RDD Lineage



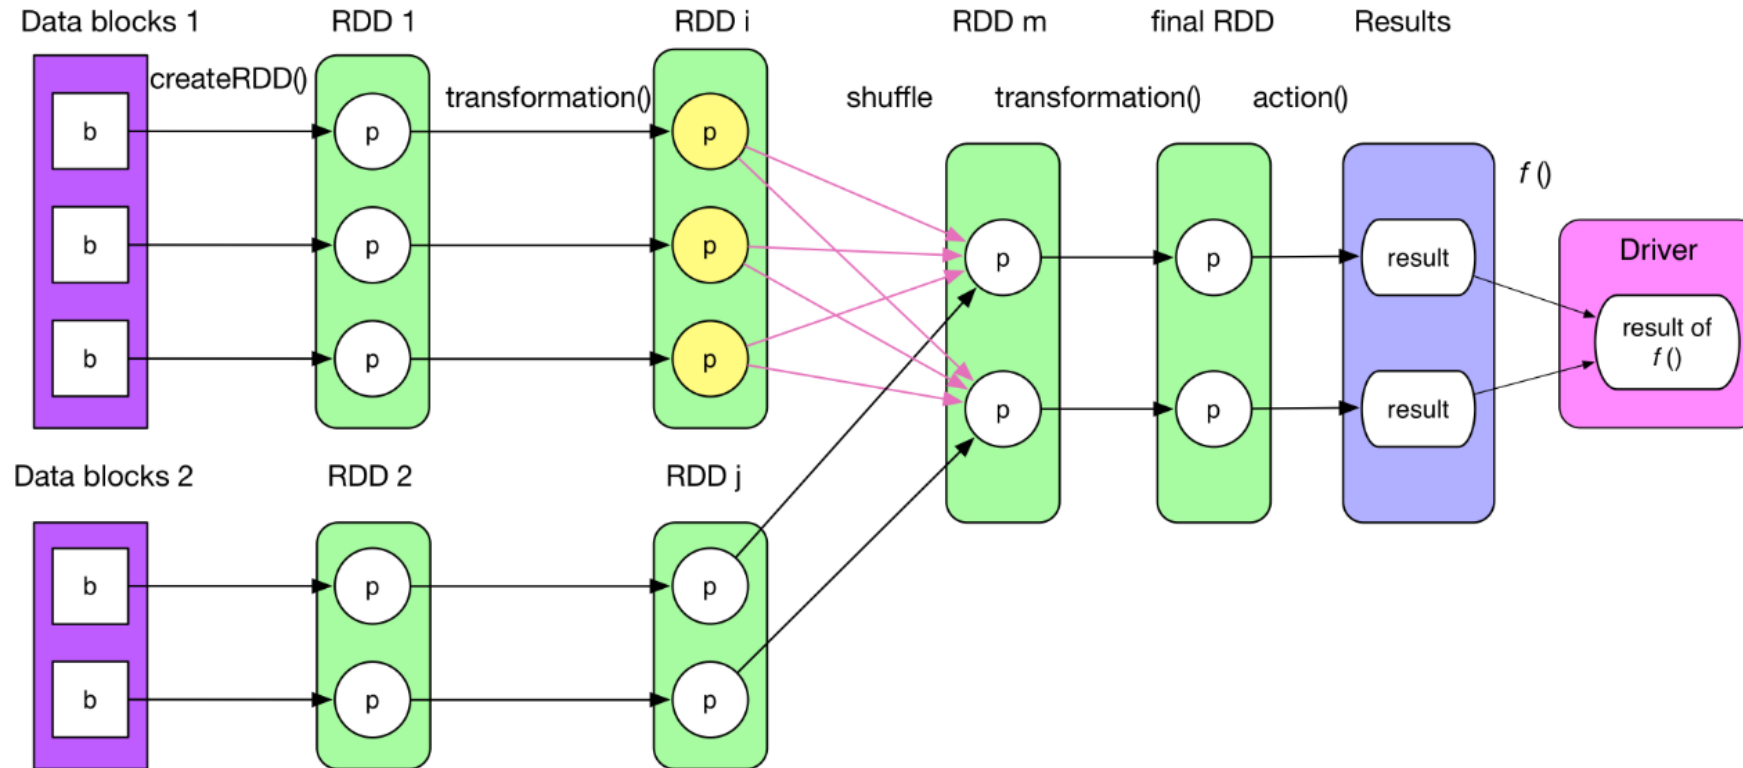
# toDebugString()

---

```
temp = graph.join(temp)
    .filter(x => x._2._2 != -1)
    .flatMap(x => x._2._1
        .map(y => (y, x._2._2 + 1)))
```

```
(17) ShuffledRDD[179] at reduceByKey at countDistance.scala:90 []
+-(17) UnionRDD[178] at union at countDistance.scala:90 []
    | MapPartitionsRDD[177] at flatMap at countDistance.scala:86 []
    | MapPartitionsRDD[176] at filter at countDistance.scala:85 []
    | MapPartitionsRDD[175] at join at countDistance.scala:84 []
    | MapPartitionsRDD[174] at join at countDistance.scala:84 []
```

# Transformation and action

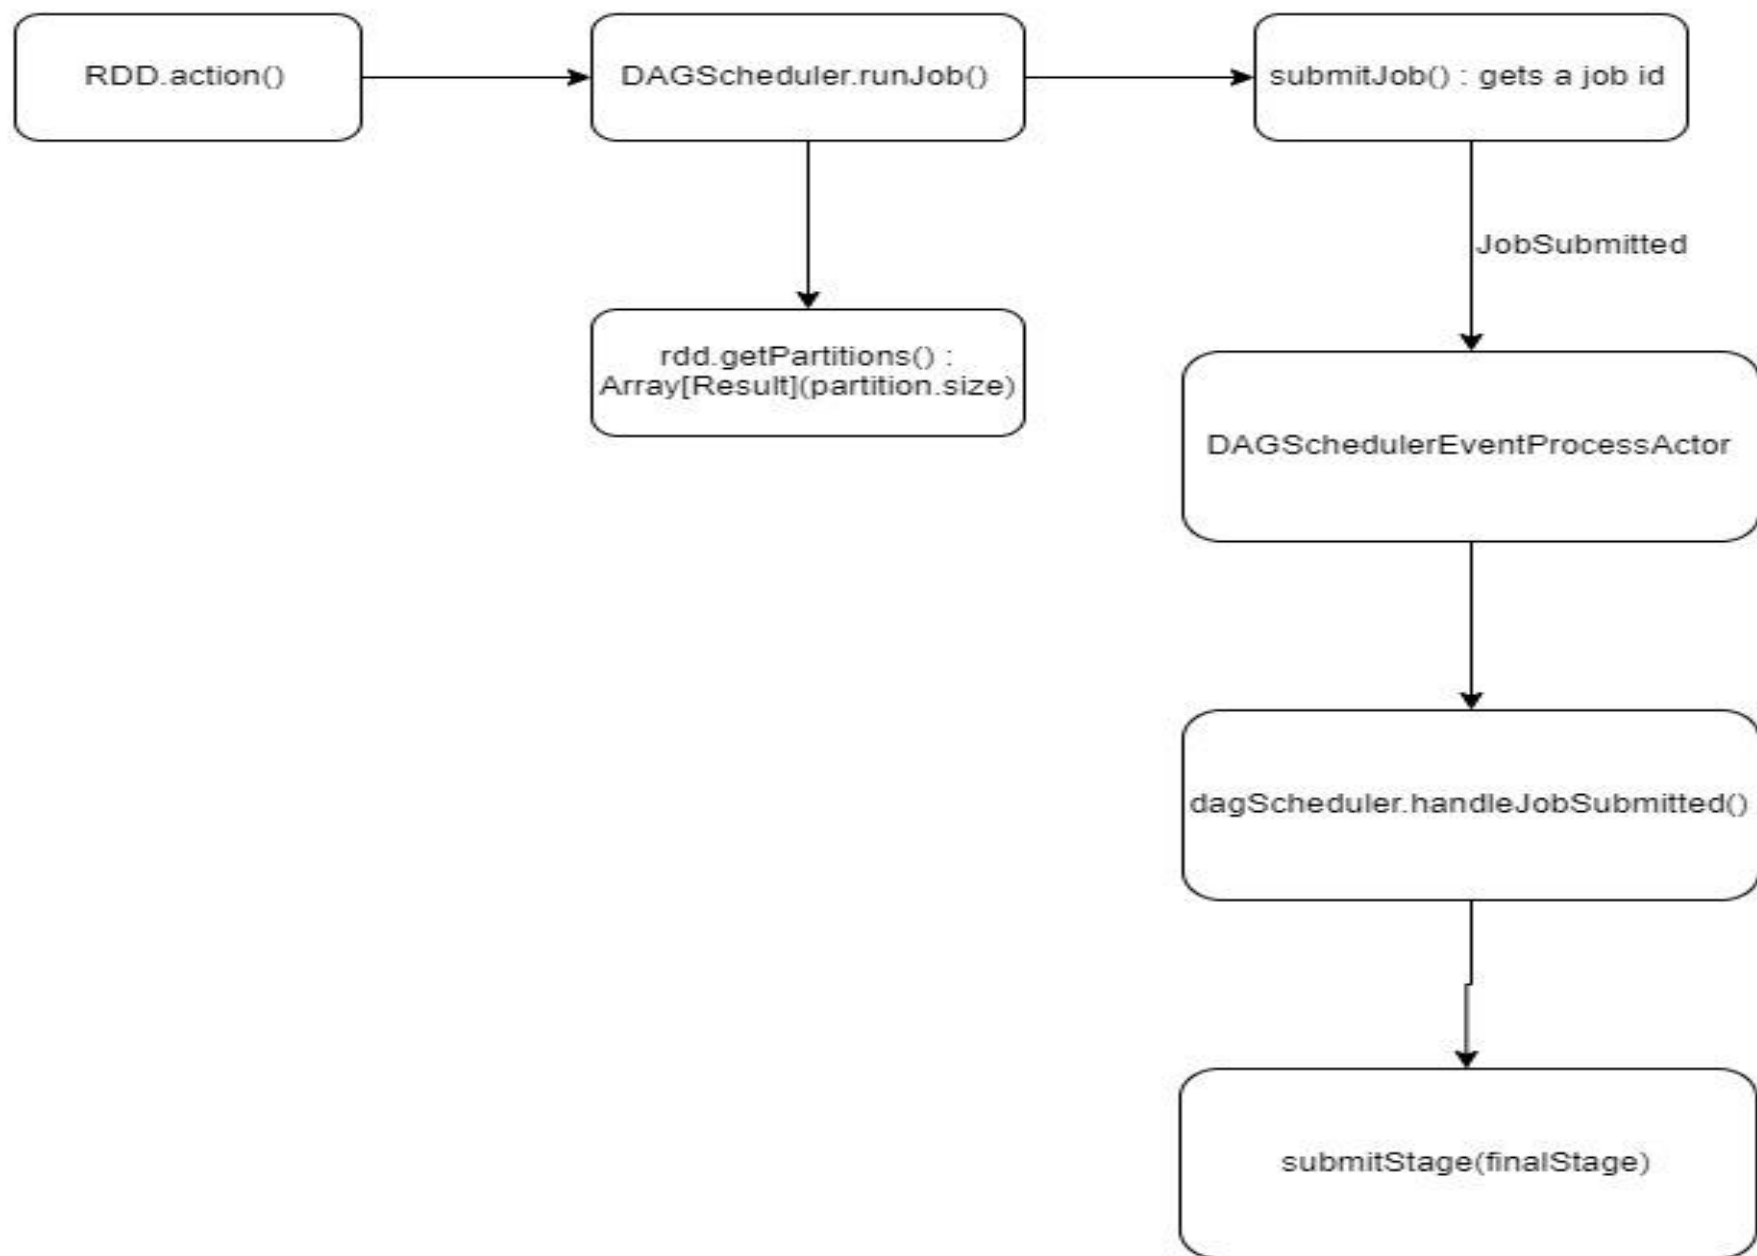


# Demo

---









# Next Sprint Goals

- Design strategies to implement N-Way merge algorithm
- Start implementing the N-Way merge algorithm.

Any Questions?

Thank You!

